
Computational Optimal Transport: LP, Quadratic, and Entropic Regularizations

Jiayi Chen

Columbia UNI: JC6601

jc6601@columbia.edu

Abstract

We study the discrete optimal transport (OT) problem on random 2D point clouds and compare three solver paradigms for the linear-programming formulation (simplex, interior point, primal–dual hybrid gradient), two algorithms for the quadratically regularized OT (ADMM and IPM), and the Sinkhorn matrix-scaling algorithm for the entropy-regularized OT. We verify all regularized variants converge to the LP optimum W_{LP} as the regularization vanishes, and we compare the discrete OT cost with the closed-form Bures–Wasserstein W_2^2 for two 2D Gaussians. Reproduction instructions are in the README of the accompanying repository.

1 Setup and notation

Let $\{x_i\}_{i=1}^n \sim \mathcal{N}(\mu_1, \Sigma_1)$ and $\{y_j\}_{j=1}^n \sim \mathcal{N}(\mu_2, \Sigma_2)$ with $\mu_1 = (0, 0)$, $\mu_2 = (4, 3)$, $\Sigma_1 = \begin{bmatrix} 1 & 0.6 \\ 0.6 & 1.2 \end{bmatrix}$, $\Sigma_2 = \begin{bmatrix} 1.5 & -0.7 \\ -0.7 & 1 \end{bmatrix}$. Both covariances are non-diagonal; uniform weights $a = b = \frac{1}{n}\mathbf{1}$; quadratic cost $C_{ij} = \|x_i - y_j\|_2^2$. We fix the random seed for reproducibility. The transportation polytope is $U(a, b) = \{P \in \mathbb{R}_{\geq 0}^{n \times n} : P\mathbf{1} = a, P^\top \mathbf{1} = b\}$.

2 Part 1: Three LP solver paradigms

We solve $W_{\text{LP}} = \min_{P \in U(a, b)} \langle C, P \rangle$ in standard form (vectorize P , sparse marginal matrix in $\mathbb{R}^{(2n) \times n^2}$). Solvers used: **Simplex**=`scipy.optimize.linprog` with HiGHS-DS, **IPM**=HiGHS-IPM, **PDLP**=Google’s PDLP via `ortools.pdlp.python` (PDHG with adaptive step sizes, preconditioning and restarts [1, 2]). We verify correctness on $n \in \{3, 4\}$: simplex and IPM agree to machine precision.

Optional: a from-scratch PDHG. We implemented the saddle-point PDHG of §4 of the handout in NumPy with $\tau = \sigma = \frac{0.99}{\sqrt{m+n}}$ (so $\tau\sigma\|K\|^2 < 1$ for $K(P) = (P\mathbf{1}, P^\top \mathbf{1})$). For the Frobenius/ ℓ_2 operator norm, $\|K\|^2 = m + n$, attained by a constant matrix. Two modifications were needed: rescale $C \leftarrow C/\|C\|_\infty$ so primal/dual scales match, and restart from the running average every 2000 steps—the lessons of PDLP [1, 4]. We then reach $\|P\mathbf{1} - a\|_1 + \|P^\top \mathbf{1} - b\|_1 < 10^{-5}$ in 13 200 iterations at $n = 200$ (Fig. 2).

Results. Table 1 summarizes wall time and accuracy at $n \in \{50, 100, 200, 500\}$ on a single CPU; Fig. 1 plots time vs. n on log–log axes. On CPU, IPM (HiGHS) is fastest except for an essentially tied $n = 50$ case, and scales roughly as $O(n^{2.5})$ in our range. Simplex (HiGHS-DS) catches IPM at small n but explodes at $n = 500$ (25 s), reflecting the combinatorial worst case of pivoting. PDLP, designed for very large LPs, is competitive (8 s at $n = 500$) but not best on CPU—its real strength is GPU vectorisation of the per-iteration work. Our custom NumPy PDHG is uniformly slower because we lack PDLP’s adaptive step sizes and Ruiz preconditioning. Critically, all three high-accuracy solvers

Table 1: Part 1 – LP solvers on random Gaussian point clouds (seed 0). Marginal violation is $\|P\mathbf{1} - a\|_1 + \|P^\top \mathbf{1} - b\|_1$; cost gap is $\langle C, P \rangle - W_{\text{LP}}$ where W_{LP} is the simplex/IPM optimum (which agree to machine precision). The reference optima for $n = 50, 100, 200, 500$ are 26.99, 25.48, 25.06, and 24.49.

n	Wall time (s)				Marginal violation				Cost gap to W_{LP}			
	50	100	200	500	50	100	200	500	50	100	200	500
Simplex (HiGHS-DS)	0.010	0.080	0.96	25.03	10^{-16}	10^{-16}	10^{-16}	10^{-16}	0	0	0	0
IPM (HiGHS-IPM)	0.011	0.054	0.25	2.44	10^{-16}	10^{-16}	10^{-16}	10^{-16}	0	0	0	0
PDLP (OR-Tools)	0.022	0.151	0.68	8.02	$3 \cdot 10^{-14}$	$6 \cdot 10^{-14}$	$5 \cdot 10^{-11}$	$6 \cdot 10^{-13}$	10^{-13}	10^{-13}	$8 \cdot 10^{-10}$	$4 \cdot 10^{-8}$
PDHG (custom NumPy)	0.092	0.357	1.52	33.72	$4 \cdot 10^{-6}$	10^{-5}	10^{-5}	10^{-4}	$2 \cdot 10^{-4}$	10^{-4}	10^{-4}	$3 \cdot 10^{-4}$

(Simplex, IPM, PDLP) agree on W_{LP} to 10^{-10} on small/medium instances; PDLP’s first-order terminal accuracy degrades mildly to a $4 \cdot 10^{-8}$ cost gap at $n = 500$. Our PDHG agrees to $\sim 10^{-5}$.

3 Part 2: Quadratic regularization

We solve $\min_{P \in U(a,b)} \langle C, P \rangle + \frac{\lambda}{2} \|P\|_F^2$. **ADMM via OSQP** (CVXPY) and **IPM via Clarabel** (CVXPY) gave costs 25.067388 and 25.064712 respectively at $\lambda = 0.1$, $n = 200$, with wall times 6.07 s (OSQP) vs. 0.31 s (Clarabel). The IPM is twenty times faster on this problem because the QP is small, strictly convex, and well-conditioned: Newton steps converge in ~ 15 iterations. ADMM via OSQP requires many more iterations and pays for the upfront KKT factorization.

Custom ADMM and the role of ρ . We also implemented the splitting from the handout in NumPy with one modification that drastically improved convergence: we put the nonnegativity constraint in the P -block (so the P -update has the closed-form clip given in eq. (4) of the handout) and put the *affine* marginal constraint in the Q -block (so the Q -update becomes the closed-form orthogonal projection onto $\{Q : Q\mathbf{1} = a, Q^\top \mathbf{1} = b\}$, no Dykstra inner loop required). Both blocks then have $O(n^2)$ closed-form updates. Fig. 5 shows the achieved marginal violation after 20 000 iterations as a function of ρ at $\lambda = 0.1$, $n = 200$: convergence requires $\rho \gtrsim n$ (e.g. $\rho = 300$ suffices), because a and b scale as $1/n$ while C has $O(1)$ entries, so the dual updates need to be amplified to balance. This matches the well-known sensitivity of vanilla ADMM to the penalty ρ and motivates the adaptive schemes used by OSQP.

Regularization path. Fig. 3 plots the gap $\langle C, P_\lambda^* \rangle - W_{\text{LP}}$ on log–log axes for $\lambda \in [10^{-3}, 10]$. The gap shrinks monotonically and is tiny ($< 10^{-2}$) over the full range; this is because $\|P^*\|_F^2 \approx 1/n$ for the sparse LP solution, so $\frac{\lambda}{2} \|P\|_F^2$ is at most ~ 0.025 even at $\lambda = 10$. Unlike entropic regularization, the quadratic penalty still admits exact zeros: the KKT condition reduces to the soft-thresholding $P_{ij}^* = \max(0, (f_i + g_j - C_{ij})/\lambda)$, so entries with $f_i + g_j < C_{ij}$ are pinned to zero (Fig. 7, center-left).

4 Part 3: Entropic regularization (Sinkhorn)

We solve $\min_{P \in U(a,b)} \langle C, P \rangle + \lambda \sum_{ij} P_{ij} \log P_{ij}$ with three implementations: (i) the POT library (`ot.sinkhorn`, log-domain), (ii) our own log-domain Sinkhorn in NumPy, and (iii) a generic conic solver through CVXPY using `cp.entr` (the exponential cone) [3, 5]. The two candidate conic backends play different algorithmic roles: Clarabel is a true interior-point conic solver (analogous to HiGHS-IPM in Part 1, but for the exponential cone), while SCS is a first-order operator-splitting conic solver, closer in spirit to ADMM than to an IPM. We work in the log domain throughout: at $\lambda = 0.01$ the entries $\exp(-C/\lambda)$ underflow at double precision, so naive matrix-form Sinkhorn fails. At $n = 200$ our custom Sinkhorn matches POT to 10^{-5} in cost; convergence is fast at $\lambda = 1$ (33 iterations to reach 10^{-6} marginal violation) and slows as λ decreases (318 iterations at $\lambda = 0.1$; at $\lambda = 0.01$, 20 000 iterations still leave a residual of $1.7 \cdot 10^{-6}$, Fig. 6).

Sinkhorn vs. conic IPM (Clarabel) and FO conic (SCS), $\lambda = 0.1$. Table 2 summarizes the three solvers across n on the same seed-0 instances. Wall times are medians of three timed repeats after one warm-up solve, and include library/Python overhead. The small cases show that conic solvers can be

competitive: Clarabel is fastest at $n=30$ and $n=100$. Sinkhorn’s time is not monotone in n because the number of scaling iterations is strongly instance-dependent; the $n=100$ instance is harder than the $n=200$ instance at this λ (22 283 iterations vs. 318). At $n=200$, however, Clarabel fails on the 40 000-entry exponential cone and SCS is roughly $60\times$ slower than Sinkhorn while giving the same transport cost to about $5\cdot 10^{-3}$.

Table 2: Part 3 – entropic OT at $\lambda=0.1$: custom log-domain Sinkhorn vs. Clarabel (interior-point conic) vs. SCS (first-order operator-splitting conic). Timings are median wall times over three post-warm-up repeats on the same seed-0 instance; the middle block shows Sinkhorn scaling iterations to reach 10^{-6} marginal violation.

n	Wall time (s)				Sinkhorn iterations				Cost $\langle C, P^* \rangle$			
	30	50	100	200	30	50	100	200	30	50	100	200
Sinkhorn (custom log)	0.259	0.028	3.20	0.161	8931	586	22283	318	24.831	27.037	25.534	25.132
Clarabel (IPM conic)	0.021	0.059	0.30	FAIL	—	—	—	—	24.831	27.037	25.534	—
SCS (FO conic)	0.120	0.258	1.54	10.27	—	—	—	—	24.832	27.044	25.543	25.138

Regularization path. Fig. 4 shows the cost gap to W_{LP} on log–log axes for $\lambda \in [10^{-2}, 10]$. The gap closes smoothly to $\sim 10^{-3}$ as $\lambda \rightarrow 0$. Note that every Sinkhorn plan is *mathematically dense*: $P_{ij}^* = u_i K_{ij} v_j > 0$ for all (i, j) since u, v are strictly positive at convergence. The edge counts annotated under Fig. 7 are therefore only the number of entries above the visualization threshold $10^{-3} P_{\max}$, not the support of P^* . By contrast the LP plan is a vertex of $U(a, b)$ and so has at most $2n-1$ structurally nonzero entries, and the quadratic plan inherits exact zeros from the clip $\max(0, \cdot)$ in the soft-thresholding KKT step; both are genuinely sparse. Thus at $\lambda=1$ the Sinkhorn plan has $\sim 25\,500$ entries above threshold (effectively dense for plotting), whereas at $\lambda=0.01$ only $\sim 1\,100$ entries remain above threshold and the plan visually approximates the LP support.

5 Part 4: Gaussian OT and Bures–Wasserstein

For the chosen μ_i, Σ_i the Bures–Wasserstein closed form $W_2^2 = \|\mu_1 - \mu_2\|^2 + \text{tr}(\Sigma_1 + \Sigma_2 - 2(\Sigma_1^{1/2} \Sigma_2 \Sigma_1^{1/2})^{1/2})$ evaluates to 25.852. We sample $n \in \{50, 100, 200, 500\}$ points from each Gaussian and solve the discrete OT LP for 10 random seeds per n (Fig. 8). Across our finite-sample range $n \leq 500$, the empirical discrete OT cost fluctuates around 25.32–25.40, i.e. slightly below the population value 25.852. We interpret this as finite-sample noise rather than a guaranteed downward bias: the sign and magnitude of the bias in plug-in empirical OT depend on the distributions and the regime, and in general one cannot assert that $\mathbb{E} \langle C, P_n^* \rangle$ lies below W_2^2 for all n (a familiar counterexample is two equal distributions, where $W_2^2 = 0$ but $\mathbb{E} \langle C, P_n^* \rangle > 0$). Asymptotically the plug-in estimator does converge to W_2^2 [5]; the trend toward the dashed line in Fig. 8 is not visually conclusive at our sample sizes, though the standard deviation does shrink with n ($1.85 \rightarrow 0.59$). Showing convergence cleanly would require substantially larger n and more seeds. Fig. 9 overlays the LP coupling at $n=200$ on the $2\text{-}\sigma$ Gaussian contour ellipses.

LLM usage and reproduction. A coding LLM was used as a pair-programmer for boilerplate (plot styling, sparse-matrix assembly, L^AT_EX scaffolding) and surfaced a bug in the first ADMM splitting; all algorithmic decisions, parameter sweeps, and analysis are the author’s, and every reported number comes from running the repository. Reproduction: `pip install -r requirements.txt && python run_all.py`.

Repository: <https://github.com/marcochen0720/EEOR6616Project>.

References

- [1] David Applegate, Mateo Díaz, Oliver Hinder, Haihao Lu, Miles Lubin, Brendan O’Donoghue, and Warren Schudy. Practical large-scale linear programming using primal-dual hybrid gradient. *Advances in Neural Information Processing Systems*, 2021.
- [2] Antonin Chambolle and Thomas Pock. A first-order primal-dual algorithm for convex problems with applications to imaging. *Journal of Mathematical Imaging and Vision*, 40(1):120–145, 2011.
- [3] Marco Cuturi. Sinkhorn distances: Lightspeed computation of optimal transport. *Advances in Neural Information Processing Systems*, 2013.
- [4] Haihao Lu and Jinwen Yang. An overview of gpu-based first-order methods for linear programming and extensions. *arXiv preprint arXiv:2503.09566*, 2025.
- [5] Gabriel Peyré and Marco Cuturi. Computational optimal transport. *Foundations and Trends in Machine Learning*, 11(5–6), 2019.

Appendix: Figures




Figure 1: Part 1: log-log wall time vs. problem size for the OT LP. IPM (HiGHS) wins on CPU; simplex blows up at $n = 500$; PDLP and our custom NumPy PDHG sit between the two.




Figure 2: Custom PDHG on the OT LP at $n = 200$: the averaged iterate $\bar{P}$ converges to $W_{LP} = 25.06$ (left); the marginal violation reaches 10^{-6} in $\sim 13\,000$ iterations (right).




Figure 3: Quadratic regularization path: cost gap to W_{LP} on log-log axes. The gap is small over the entire range $\lambda \in [10^{-3}, 10]$ because $\|P^*\|_F^2 \approx 1/n$ at the LP optimum.




Figure 4: Sinkhorn regularization path: cost gap to W_{LP} on log-log axes. Smooth $O(\lambda)$ behaviour visible across two decades.




Figure 5: Effect of the ADMM penalty ρ in our custom NumPy ADMM (quadratic OT, $\lambda = 0.1$, $n = 200$). Convergence requires $\rho \gtrsim n$ (here $\rho \geq 300$); below that, 20 000 iterations are insufficient and the marginal violation stalls.




Figure 6: Sinkhorn convergence at $n = 200$ for $\lambda \in \{1, 0.1, 0.01\}$. Linear convergence in iterates; rate degrades as $\lambda \rightarrow 0$.




Figure 7: Same point cloud ($n = 200$, seed 0) under four solvers. Top row: log-magnitude heatmap of the coupling $P/P_{\max}$ with source/target indices sorted by x -coordinate. Bottom row: source-target connections weighted by P_{ij} . The number of entries *above the visualization threshold* $10^{-3}P_{\max}$ is annotated below each panel. The LP and quadratic ($\lambda = 0.1$) plans are genuinely sparse (LP at a vertex of $U(a, b)$; quadratic gets exact zeros from the soft-thresholding KKT step). Both Sinkhorn plans are mathematically dense ($P_{ij}^* > 0$ everywhere); the apparent sparsity at $\lambda = 0.01$ (≈ 1100 entries above threshold) is numerical, not structural.




Figure 8: Part 4: discrete OT cost vs. n , mean $\pm$ one standard deviation across 10 seeds. Dashed line is the Bures–Wasserstein closed form $W_2^2 = 25.85$. In this finite-sample range the empirical estimates fluctuate slightly below the population value; this should be interpreted as finite-sample noise rather than a universal downward bias. Standard deviation shrinks with n .




Figure 9: Part 4: LP coupling at $n = 200$ between samples of two 2D Gaussians, overlaid on the 2σ contour ellipses (blue source, red target). Arrow opacity/width is proportional to P_{ij} .